

# **Model Selection, Evaluation, Diagnosis**

INFO-4604, Applied Machine Learning  
University of Colorado Boulder

**November 8, 2020**

Prof. Abe Handler

A little bit of review

# Evaluation

In homework, you've seen that:

- training data is usually separate from **test data**
- training accuracy is often much higher than test accuracy
  - Training accuracy is what your classifier is optimizing for (plus regularization), but not a good indicator of how it will perform

# Evaluation


Distinction between:

- **in-sample** data
  - The data that is available when building your model
  - “Training” data in machine learning terminology
- **out-of-sample** data
  - Data that was not seen during training
  - Also called **held-out data** or a **holdout set**
  - Useful to see what your classifier will do on data it hasn't seen before
  - Usually assumed to be from the same distribution as in-sample data

# Evaluation

- Ideally, you should be “blind” to the test data until you are ready to evaluate your final model
- Often you need to evaluate a model repeatedly (e.g., you’re trying to pick the best regularization strength, and you want to see how different values affect the performance)
- So you set aside part of training data as a validation set, or perform cross-validation

# Evaluation



# Cross validation

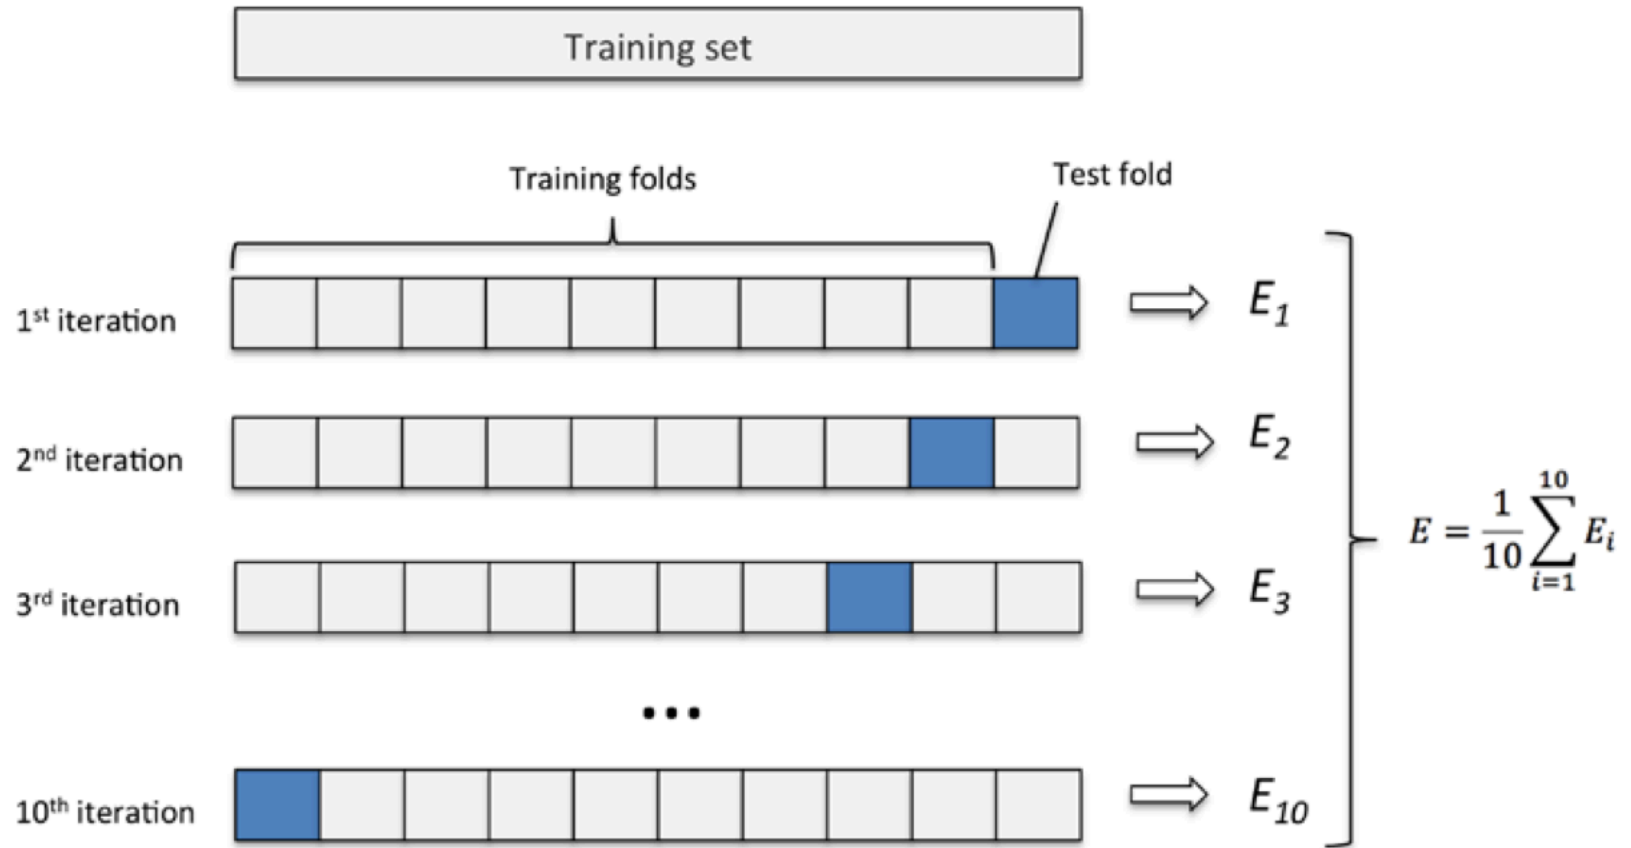


Illustration of 10-fold cross-validation

# Splitting your dataset

- Remember to watch for leaks between train in and test. If you already know that person A likes bananas from train, you might be able to infer that person A also likes apples in test.
- This takes experience and thought. Start practicing!



# Other considerations

- If test performance looks too good to be true, it probably is.
- Watch out for bugs and dumb errors.
- This is very common everywhere, including in fancy CS research.
- Because ML algorithms are often stochastic, debugging can be trickier

# Evaluation Metrics

How do we measure performance?

What *metrics* should be used?

# Evaluation Metrics

So far, we have mostly talked about **accuracy** in this class

- The number of correctly classified instances divided by the total number of instances

**Error** is the complement of accuracy

- $\text{Accuracy} = 1 - \text{Error}$
- $\text{Error} = 1 - \text{Accuracy}$

# Evaluation Metrics

Accuracy/error give an overall summary of model performance, though sometimes hard to interpret

Example: fraud detection in bank transactions

- 99.9% of instances are legitimate
- A classifier that never predicts fraud would have an accuracy of 99.9%
- Need a better way to understand performance

# Evaluation Metrics

Some metrics measure performance with respect to a particular class

With respect to a class  $c$ , we define a prediction as:

- True positive: the label is  $c$  and the classifier predicted  $c$
- False positive: the label is not  $c$  but the classifier predicted  $c$
- True negative: the label is not  $c$  and the classifier did not predict  $c$
- False negative: the label is  $c$  but the classifier did not predict  $c$

# Evaluation Metrics

Two different types of errors:

- False positive (“type I” error)
- False negative (“type II” error)

Usually there is a tradeoff between these two

- Can optimize for one at the expense of the other
- Which one to favor? Depends on task

# Evaluation Metrics

**Precision** is the percentage of instances predicted to be positive that were actually positive

$$PRE = \frac{TP}{TP + FP}$$

Fraud example:

- Low precision means you are classifying legitimate transactions as fraudulent

# Evaluation Metrics

**Recall** is the percentage of positive instances that were predicted to be positive

$$REC = TPR = \frac{TP}{P} = \frac{TP}{FN + TP}$$

Fraud example:

- Low recall means there are fraudulent transactions that you aren't detecting



# Precision and recall

- Today, just focus on understanding what they are
- Later this week we will do some whiteboard exercises to build intuition for these numbers
- There is a deep, deep tradeoff between making wrong predictions and missing things.
- Ideally, you want as many of both as possible. But in reality you have to choose
- Good practitioners should be able to help non-technical people understand and make this choice

# Evaluation Metrics

Similar to optimizing for false positives vs false negatives, there is usually a tradeoff between prediction and recall

- Can increase one at the expense of the other
- One might be more important than the other, or they might be equally important; depends on task

Fraud example:

- If a human is reviewing the transactions flagged as fraudulent, probably optimize for recall
- If the classifications are taken as-is (no human review), probably optimize for precision

# Evaluation Metrics

Can modify prediction rule to adjust tradeoff

- Increased prediction threshold (i.e., score or probability of an instance belonging to a class)
  - increased precision
    - Fewer instances will be predicted positive
    - But the ones that are classified positive are more likely to be classified correctly (more confidence classifier)
- Decreased threshold → increased recall
  - More instances will get classified as positive (the bar has been lowered)
  - But this will make your classifications less accurate, lower precision

# Evaluation Metrics

The **F1 score** is an average of precision and recall

- Used as a summary of performance, still with respect to a particular class  $c$
- Defined using *harmonic* mean, affected more by lower number
  - Both numbers have to be high for F1 to be high
  - F1 is therefore useful when both are important

$$F1 = 2 \frac{PRE \times REC}{PRE + REC}$$

# Evaluation Metrics

## Which metrics to use?

- Accuracy easier to understand/communicate than precision/recall/F1, but harder to interpret correctly
- Precision/recall/F1 generally more informative
  - If you have a small number of classes, just show P/R/F for all classes
  - If you have a large number of classes, then probably should do macro/micro averaging
- F1 better if precision/recall both important, but sometimes you might highlight one over the other

# Evaluation Metrics

It is often a good idea to contextualize your results by comparing to a *baseline* level of performance

- A “dummy” baseline, like always outputting the majority class, can be useful to understand if your data is imbalanced (like in fraud example)
- A simple, “default” classifier for your task, like using standard 1-gram features for text, can help you understand if your modifications resulted in an improvement

# Model Selection

Evaluation can help you choose (or *select*) between competing models

- Which classification algorithm to use?
- Best preprocessing steps?
- Best feature set?
- Best hyperparameters?

Usually these are all decided *empirically* based on testing different possibilities on your data

# Model Selection

Selecting your model to get optimal test performance is risky

- What works best for the particular test set might not actually be the best in general
- Overfitting to the test data

**Validation** (or **development**) data refers to data that is held-out for measuring performance, but is separate from the final test set



# Model Selection

What settings should you test? Lots of possibilities.

Different decisions/hyperparameters depend on each other; not independent

- e.g., optimal regularization strength depends on what kind of regularizer (L2 vs L1), what kind of feature selection, etc

Best to optimize *combinations* of settings, rather than optimizing individually

# Model Selection

A **grid search** is the process of evaluating every combination of settings (from a specified set of potential values) on validation data

Quickly becomes expensive... example:

- 5 regularization values
- 2 regularizer types
- 3 kernel settings
- 5 feature selection settings
- 2 preprocessing options
- = 300 combinations

# Model Selection

Might only perform a grid search on a subset of combinations initially

One option: start by searching a small number of very different values (e.g.,  $C=\{0.1, 1.0, 10.0, 100.0\}$ ), then fine-tune with more values close to the optimal one (e.g. find that  $C=1.0$  and  $C=10.0$  are the best of the above options, so now try  $C=\{1.0, 2.0, 4.0, 6.0, 8.0, \dots\}$ )

# Understanding and Diagnosis

In addition to *measuring* performance, also important to *understand* performance

Want to be able to answer:

- Why does the model make the predictions it does?
- What kinds of errors does the model make?
- What can be done to improve the model?

# Error Analysis

What kinds of mistakes does a model make?

Which classes does a classifier tend to mix up?

# Error Analysis



# Error Analysis

A **confusion matrix** (or **error matrix**) is a table that counts the number of test instances with each true label vs each predicted label.

		Predicted			
		Iris-setosa	Iris-versicolor	Iris-virginica	$\Sigma$
Actual	Iris-setosa	100.0 %	0.0 %	0.0 %	50
	Iris-versicolor	0.0 %	88.7 %	6.4 %	50
	Iris-virginica	0.0 %	11.3 %	93.6 %	50
$\Sigma$		50	53	47	150

From: <https://docs.orange.biolab.si/3/visual-programming/widgets/evaluation/confusionmatrix.html>

# Error Analysis

A **confusion matrix** (or **error matrix**) is a table that counts the number of test instances with each true label vs each predicted label.

- In binary classification, the confusion matrix is just a 2x2 table of true/false positive/negatives
- But can be generalized to multiclass settings



# Error Analysis

Just as false negatives vs false positives are two different types of errors where one may be preferable, different types of multiclass errors may have different importance

- Mistaking a deer for an antelope is not such a bad mistake
- Mistaking a deer for a cereal box would be an odd mistake

Are the mistakes acceptable? Need to look at confusions, not just a summary statistic.

# Error Analysis

Another way to understand why your classifier is behaving a certain way is to examine the parameters that are learned after training

- e.g., the decision tree structure, or the weight vector values

If features are associated with classes in a way that doesn't make sense to you, that might mean the model is not working the way you intended

- Caveat: features interact in ways that can be hard to understand, so unintuitive parameters not necessarily wrong

# Error Analysis

Another good practice: look at a sample of misclassified instances

test11.png



**msft\_captions**

a close up of a stuffed animal (55)

**msft\_tags**

indoor (94) , food (87) , bread (74) , dessert (30)

From: <https://medium.freecodecamp.org/chihuahua-or-muffin-my-search-for-the-best-computer-vision-api-cbda4d6b425d>

# Error Analysis

Another good practice: look at a sample of misclassified instances

- Might help you understand why the classifier is making those mistakes
- Might help you understand what kinds of instances the classifier makes mistakes on
  - Maybe they were ambiguous to begin with, so not surprising the classifier had trouble

# Error Analysis

Another good practice: look at a sample of misclassified instances

- If you have multiple models, can compare how they do on individual instances
- If your model outputs probabilities, useful to examine
  - If the correct class was the 2nd most probable, that's a better mistake than if it was the 10th most probable

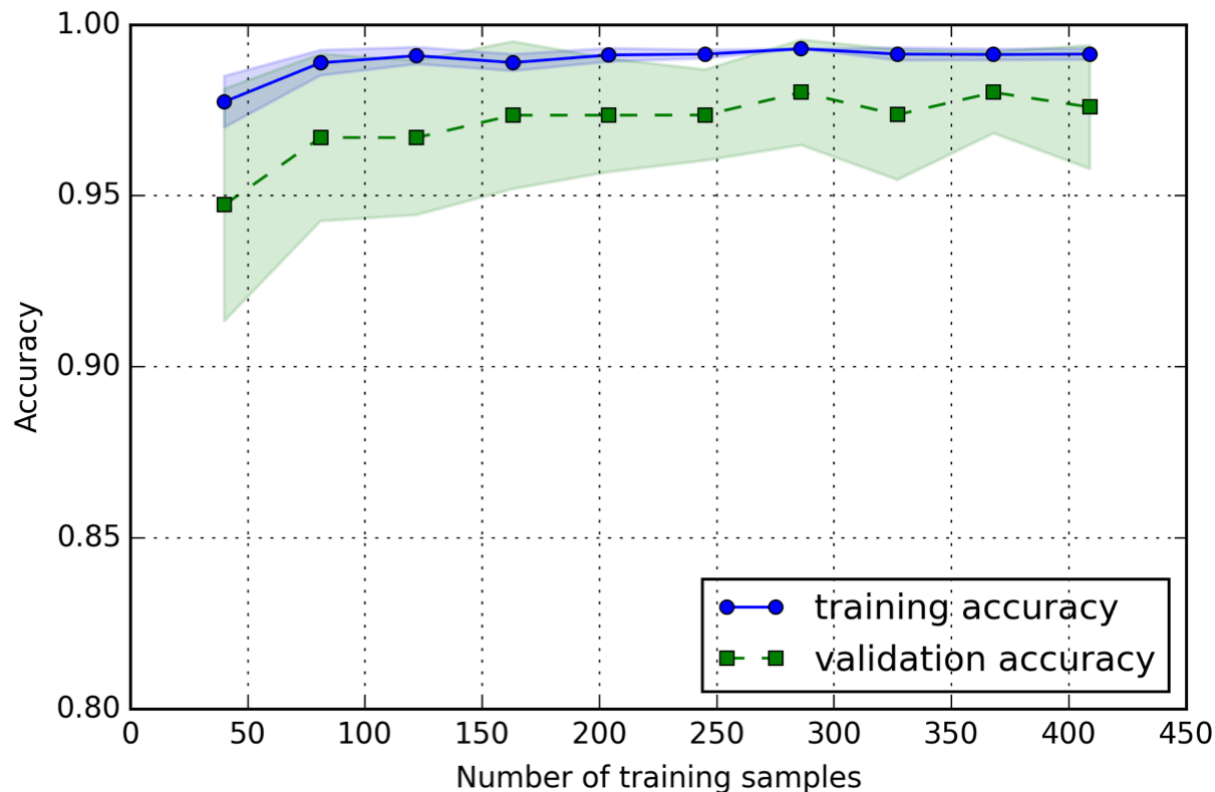
# Error Analysis

Error analysis can help inform:

- Feature engineering
  - If you observe that certain classes are more easily confused, think about creating new features that could distinguish those classes
- Feature selection
  - If you observe that certain features might be hurting performance (maybe the classifier is picking up on an association between a feature and a class that isn't meaningful), you could remove it

# Learning Curves

A **learning curve** measures the performance of a model at different amounts of training data



# Learning Curves

Primarily used to understand two things:

- How much training data is needed?
- Bias/variance tradeoff



# Learning Curves

How much training data is needed?

Usually the validation accuracy increases noticeable after an initial increase in training data, then levels off after a while

- Might still be increasing, but with diminishing returns

The learning curve can be used to predict the benefit you'll get from adding more training data

- Still increasing rapidly → get more data
- Completely flattened out → more data won't help
- Gradually increasing → need a lot more data to help

# Learning Curves


A typical pattern is that:

- Training accuracy decreases with more data
- Validation accuracy increases
- The two accuracies should converge to be similar


Something to look for:

- If gap between train/validation performance isn't closing, probably too much variance (overfitting)
- If gap between train/validation closes quickly, might suggest high bias (underfitting)


**High bias**



**High variance**




**Good bias-variance trade-off**



- Training accuracy
- Validation accuracy
- Desired accuracy

# Validation Curves

A **validation curve** measures the performance of a model at different hyperparameter settings




# Validation Curves

Validation curves help you understand the effect of hyperparameters, and also help understand the bias/variance tradeoff


- Want to find a setting where train/validation performance is similar (low variance)
- Of the settings where train/validation performance are similar, pick the one with the highest validation accuracy (low bias)

# Precision-Recall Curve



From: <https://stackoverflow.com/questions/33294574/good-roc-curve-but-poor-precision-recall-curve>

# ROC Curve



From: <https://stackoverflow.com/questions/33294574/good-roc-curve-but-poor-precision-recall-curve>

# Debugging

Lots of places in the pipeline where there could be an implementation mistake:

- Data preprocessing
- Feature engineering
- Training algorithm
- Validation pipeline

Best to start simple, compare to existing data/systems, then expand from there