

# **Optimization and Gradient Descent**

INFO-4604, Applied Machine Learning  
University of Colorado Boulder

**September 11, 2018**

Prof. Abe Handler

# Prediction Functions

Remember: a **prediction function** is the function that predicts what the output should be, given the input.

# Prediction Functions

Linear regression:

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

Linear classification (perceptron):

$$f(\mathbf{x}) = \begin{cases} 1, & \mathbf{w}^T \mathbf{x} + b \geq 0 \\ -1, & \mathbf{w}^T \mathbf{x} + b < 0 \end{cases}$$

Need to *learn* what  $\mathbf{w}$  should be!

# Learning Parameters

Goal is to learn to minimize error

- Ideally: true error
- Instead: training error

The **loss function** gives the training error when using parameters  $\mathbf{w}$ , denoted  $L(\mathbf{w})$ .

- Also called **cost function**
- More general: **objective function**  
(in general objective could be to minimize or maximize;  
with loss/cost functions, we want to minimize)

# Learning Parameters

Goal is to minimize loss function.


How do we minimize a function?

Let's review some math.

# Rate of Change

The slope of a line is also called the **rate of change** of the line.


$$y = \frac{1}{2}x + 1$$



# Rate of Change

For nonlinear functions, the “rise over run” formula gives you the average rate of change between two points

$$f(x) = x^2$$




Average slope from  
 $x = -1$  to  $x = 0$  is:  
-1

# Rate of Change

There is also a concept of rate of change at individual points (rather than two points)

$$f(x) = x^2$$




Slope at  $x=-1$  is:  
-2



# Rate of Change

The slope at a point is called the **derivative** at that point

$$f(x) = x^2$$



Intuition:

Measure the slope between two points that are really close together


# Rate of Change

The slope at a point is called the **derivative** at that point

Intuition: Measure the slope between two points that are really close together

$$\frac{f(x + c) - f(x)}{c}$$

Limit as  $c$  goes to zero



# Maxima and Minima

Whenever there is a peak in the data, this is a **maximum**

The **global** maximum is the highest peak in the entire data set, or the largest  $f(x)$  value the function can output

A **local** maximum is any peak, when the rate of change switches from positive to negative


# Maxima and Minima

Whenever there is a trough in the data, this is a **minimum**

The **global** minimum is the lowest trough in the entire data set, or the smallest  $f(x)$  value the function can output

A **local** minimum is any trough, when the rate of change switches from negative to positive

# Maxima and Minima



From: <https://www.mathsisfun.com/algebra/functions-maxima-minima.html>

All global maxima and minima are also local maxima and minima

# Derivatives

The derivative of  $f(x) = x^2$  is  $2x$

Other ways of writing this:

$$f'(x) = 2x$$

$$d/dx [x^2] = 2x$$

$$df/dx = 2x$$


The derivative is also a function! It depends on the value of  $x$ .

- The rate of change is different at different points


# Derivatives

The derivative of  $f(x) = x^2$  is  $2x$

$f(x)$



$f'(x)$



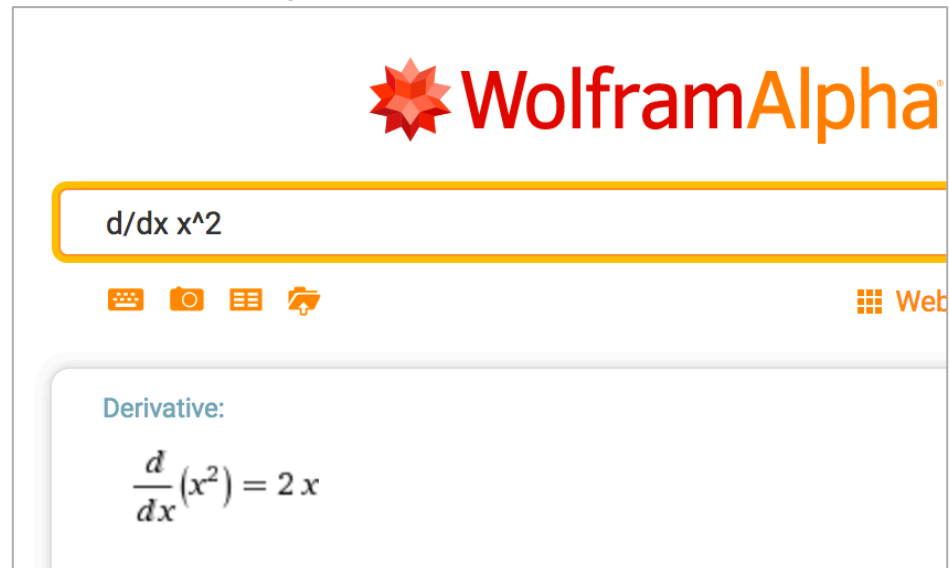
# Derivatives

How to calculate a derivative?

- Not going to do it in this class.

Some software can do it for you.

- [Wolfram Alpha](https://www.wolframalpha.com)





# Derivatives

What if a function has multiple arguments?

Ex:  $f(x_1, x_2) = 3x_1 + 5x_2$

$df/dx_1 = 3 + 5x_2$     The derivative “with respect to”  $x_1$

$df/dx_2 = 3x_1 + 5$     The derivative “with respect to”  $x_2$


These two functions are called **partial derivatives**.

The vector of all partial derivatives for a function  $f$  is called the **gradient** of the function:

$$\nabla f(x_1, x_2) = \langle df/dx_1, df/dx_2 \rangle$$

# Finding Minima

The derivative is *zero* at any local maximum or minimum.



# Finding Minima

The derivative is *zero* at any local maximum or minimum.

One way to find a minimum: set  $f'(x)=0$  and solve for  $x$ .

$$f(x) = x^2$$

$$f'(x) = 2x$$

$$f'(x) = 0 \text{ when } x = 0, \text{ so minimum at } x = 0$$

# Finding Minima

The derivative is *zero* at any local maximum or minimum.

One way to find a minimum: set  $f'(x)=0$  and solve for  $x$ .

- For most functions, there isn't a way to solve this.
- Instead: algorithmically search different values of  $x$  until you find one that results in a gradient near 0.

# Finding Minima


If the derivative is positive, the function is **increasing**.

- Don't move in that direction, because you'll be moving away from a trough.

If the derivative is negative, the function is **decreasing**.

- Keep going, since you're getting closer to a trough

# Finding Minima




$$f'(-1) = -2$$

At  $x=-1$ , the function is decreasing as  $x$  gets larger. This is what we want, so let's make  $x$  larger.

Increase  $x$  by the size of the gradient:

$$-1 + 2 = 1$$

# Finding Minima




$$f'(-1) = -2$$

At  $x=-1$ , the function is decreasing as  $x$  gets larger. This is what we want, so let's make  $x$  larger.

Increase  $x$  by the size of the gradient:

$$-1 + 2 = 1$$

# Finding Minima



$$f'(1) = 2$$


At  $x=1$ , the function is increasing as  $x$  gets larger. This is not what we want, so let's make  $x$  smaller.

Decrease  $x$  by the size of the gradient:

$$1 - 2 = -1$$



# Finding Minima




$$f'(1) = 2$$

At  $x=1$ , the function is increasing as  $x$  gets larger. This is not what we want, so let's make  $x$  smaller.

Decrease  $x$  by the size of the gradient:

$$1 - 2 = -1$$


# Finding Minima



We will keep jumping between the same two points this way.


We can fix this by using a learning rate or step size.

# Finding Minima



$$f'(-1) = -2$$
$$x \pm 2\eta =$$

# Finding Minima




$$f'(-1) = -2$$

$$x \pm 2\eta =$$

Let's use  $\eta = 0.25$ .


# Finding Minima



$$f'(-1) = -2$$

$$x = -1 + 2(.25) = -0.5$$

# Finding Minima




$$f'(-1) = -2$$

$$x = -1 + 2(.25) = -0.5$$

$$f'(-0.5) = -1$$

$$x = -0.5 + 1(.25) = -0.25$$

# Finding Minima



$$f'(-1) = -2$$

$$x = -1 + 2(.25) = -0.5$$


$$f'(-0.5) = -1$$

$$x = -0.5 + 1(.25) = -0.25$$

$$f'(-0.25) = -0.5$$

$$x = -0.25 + 0.5(.25) = -0.125$$

# Finding Minima



$$f'(-1) = -2$$

$$x = -1 + 2(.25) = -0.5$$

$$f'(-0.5) = -1$$

$$x = -0.5 + 1(.25) = -0.25$$

$$f'(-0.25) = -0.5$$

$$x = -0.25 + 0.5(.25) = -0.125$$

Eventually we'll reach  $x=0$ .



# Gradient Descent

1. Initialize the parameters  $\mathbf{w}$  to some guess (usually all zeros, or random values)
2. Update the parameters:  
$$\mathbf{w} = \mathbf{w} - \eta \nabla L(\mathbf{w})$$
3. Update the learning rate  $\eta$  (How? Later...)
4. Repeat steps 2-3 until  $\nabla L(\mathbf{w})$  is close to zero.

# Gradient Descent

Gradient descent is guaranteed to eventually find a *local* minimum if:

- the learning rate is decreased appropriately;
- a finite local minimum exists (i.e., the function doesn't keep decreasing forever).

# Gradient Ascent

What if we want to find a local *maximum*?

Same idea, but the update rule moves the parameters in the opposite direction:

$$\mathbf{w} = \mathbf{w} + \eta \nabla L(\mathbf{w})$$

# Learning Rate

In order to guarantee that the algorithm will converge, the learning rate should decrease over time. Here is a general formula.

At iteration  $t$ :

$$\eta_t = c_1 / (t^a + c_2),$$

$$\text{where } 0.5 < a < 2$$

$$c_1 > 0$$

$$c_2 \geq 0$$

# Stopping Criteria

For most functions, you probably won't get the gradient to be exactly equal to **0** in a reasonable amount of time.

Once the gradient is sufficiently close to **0**, stop trying to minimize further.

How do we measure how close a gradient is to **0**?

# Distance

A special case is the distance between a point and zero (the *origin*).

$$d(\mathbf{p}, \mathbf{0}) = \sqrt{\sum_{i=1}^k (p_i)^2}$$

This is called the **Euclidean norm** of  $\mathbf{p}$

- A norm is a measure of a vector's length
- The Euclidean norm is also called the **L2 norm**

# Distance

A special case is the distance between a point and zero (the *origin*).

$$d(\mathbf{p}, \mathbf{0}) = \sqrt{\sum_{i=1}^k (p_i)^2}$$

Also written:  $\|\mathbf{p}\|$

# Stopping Criteria

Stop when the norm of the gradient is below some threshold,  $\theta$ :

$$\|\nabla L(\mathbf{w})\| < \theta$$

Common values of  $\theta$  are around .01, but if it is taking too long, you can make the threshold larger.



# Gradient Descent

1. Initialize the parameters  $\mathbf{w}$  to some guess (usually all zeros, or random values)
2. Update the parameters:  
$$\mathbf{w} = \mathbf{w} - \eta \nabla L(\mathbf{w})$$
$$\eta = c_1 / (t^a + c_2)$$
3. Repeat step 2 until  $\|\nabla L(\mathbf{w})\| < \theta$  or until the maximum number of iterations is reached.

End Sep 11th

# Revisiting Perceptron

In perceptron, you increase the weights if they were an underestimate and decrease if they were an overestimate.

$$w_j += \eta (y_i - f(x_i)) x_{ij}$$

This looks similar to the gradient descent rule.

- Is it? We'll come back to this.

# Adaline

Similar algorithm to perceptron (but uncommon):

Predictions use the same function:

$$f(\mathbf{x}) = \begin{cases} 1, & \mathbf{w}^T \mathbf{x} \geq 0 \\ -1, & \mathbf{w}^T \mathbf{x} < 0 \end{cases}$$

(here the bias  $b$  is folded into the weight vector  $\mathbf{w}$ )

# Adaline

Perceptron minimizes the number of errors.

Adaline instead tries to make  $\mathbf{w}^T \mathbf{x}$  close to the correct value (1 or -1, even though  $\mathbf{w}^T \mathbf{x}$  can be any real number).

Loss function for Adaline:

$$L(\mathbf{w}) = \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

This is called the **squared error**.  
(This is the same loss function used for linear regression.)

# Adaline

What is the derivative of the loss?

$$L(\mathbf{w}) = \sum_{i=1}^N (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

$$dL/dw_j = \sum_{i=1}^N -2 x_{ij} (y_i - \mathbf{w}^T \mathbf{x}_i)$$

# Adaline

The gradient descent algorithm for Adaline updates each feature weight using the rule:

$$w_j += \eta \sum_{i=1}^N 2 x_{ij} (y_i - \mathbf{w}^T \mathbf{x}_i)$$

Two main differences from perceptron:

- $(y_i - \mathbf{w}^T \mathbf{x}_i)$  is a real value, instead of a binary value (perceptron either correct or incorrect)
- The update is based on the entire training set, instead of one instance at a time.

# Adaline

The gradient descent algorithm for Adaline updates each feature weight using the rule:

$$w_j += \eta \sum_{i=1}^N 2 x_{ij} (y_i - \mathbf{w}^T \mathbf{x}_i)$$

Two main differences from perceptron:

- $(y_i - \mathbf{w}^T \mathbf{x}_i)$  is a real value, instead of a binary value (perceptron either correct or incorrect)
- The update is based on the entire training set, instead of one instance at a time.



# Stochastic Gradient Descent

A variant of gradient descent makes updates using an approximate of the gradient that is only based on one instance at a time.

$$L_i(\mathbf{w}) = (y_i - \mathbf{w}^T \mathbf{x}_i)^2$$

$$dL_i/dw_j = -2 x_{ij} (y_i - \mathbf{w}^T \mathbf{x}_i)$$

# Stochastic Gradient Descent

General algorithm for SGD:

1. Iterate through the instances in a random order
  - a) For each instance  $x_i$ , update the weights based on the gradient of the loss for that instance only:

$$\mathbf{w} = \mathbf{w} - \eta \nabla L_i(\mathbf{w}; \mathbf{x}_i)$$

The gradient for one instance's loss is an approximation to the true gradient

- stochastic = random  
The *expected* gradient is the true gradient

# Adaline

The gradient descent algorithm for Adaline updates each feature weight using the rule:

$$w_j += \eta \sum_{i=1}^N 2 x_{ij} (y_i - \mathbf{w}^T \mathbf{x}_i)$$


Two main differences from perceptron:

- $(y_i - \mathbf{w}^T \mathbf{x}_i)$  is a real value, instead of a binary value (perceptron either correct or incorrect)
- The update is based on the entire training set, instead of one instance at a time.

# Revisiting Perceptron

Perceptron has a different loss function:


$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^\top \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^\top \mathbf{x}_i), & \text{otherwise} \end{cases}$$



# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^\top \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^\top \mathbf{x}_i), & \text{otherwise} \end{cases}$$




The derivative here is 0.  
No gradient descent  
updates if the prediction  
was correct.

# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^T \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^T \mathbf{x}_i), & \text{otherwise} \end{cases}$$




The derivative here is  $-y_i x_{ij}$ .  
If  $x_{ij}$  is positive,  $dL_i/w_j$  will be negative when  $y_i$  is positive, so the gradient descent update will be positive.

# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^T \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^T \mathbf{x}_i), & \text{otherwise} \end{cases}$$




The derivative here is  $-y_i x_{ij}$ . If  $x_{ij}$  is positive,  $dL_i/w_j$  will be negative when  $y_i$  is positive, so the gradient descent update will be positive.

# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^T \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^T \mathbf{x}_i), & \text{otherwise} \end{cases}$$




The derivative here is  $-y_i x_{ij}$ .  
If  $x_{ij}$  is positive,  $dL_i/w_j$  will be positive when  $y_i$  is negative, so the gradient descent update will be negative.



# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^\top \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^\top \mathbf{x}_i), & \text{otherwise} \end{cases}$$




The derivative doesn't actually exist at this point (the function isn't smooth)

# Revisiting Perceptron

Perceptron has a different loss function:

$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^\top \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^\top \mathbf{x}_i), & \text{otherwise} \end{cases}$$



A **subgradient** is a generalization of the gradient for points that are not differentiable.

0 and  $-y_i \mathbf{x}_{ij}$  are both valid subgradients at this point.

# Revisiting Perceptron

Perceptron has a different loss function:



$$L_i(\mathbf{w}; \mathbf{x}_i) = \begin{cases} 0, & y_i (\mathbf{w}^\top \mathbf{x}_i) \geq 0 \\ -y_i (\mathbf{w}^\top \mathbf{x}_i), & \text{otherwise} \end{cases}$$

Perceptron is a stochastic gradient descent algorithm using this loss function (and using the subgradient instead of gradient)

# Convexity

How do you know if you've found the global minimum, or just a local minimum?


A **convex** function has only one minimum:



# Convexity



How do you know if you've found the global minimum, or just a local minimum?

A **convex** function has only one minimum:



# Convexity

A **concave** function has only one maximum:




Sometimes people use “convex” to mean either convex or concave

# Convexity

Squared error is a convex loss function, as is the perceptron loss.

Note: convexity means there is only one minimum value, but there may be multiple parameters that result in that minimum value.



# Summary

Most machine learning algorithms are some combination of a loss function + an algorithm for finding a local minimum.

- Gradient descent is a common minimizer, but there are others.

With most of the common classification algorithms, there is only one global minimum, and gradient descent will find it.

- Most often: supervised functions are convex, unsupervised functions are non-convex.



# Summary

1. Initialize the parameters  $\mathbf{w}$  to some guess (usually all zeros, or random values)
2. Update the parameters:  
$$\mathbf{w} = \mathbf{w} - \eta \nabla L(\mathbf{w})$$
$$\eta = c_1 / (t^a + c_2)$$
3. Repeat step 2 until  $\|\nabla L(\mathbf{w})\| < \theta$  or until the maximum number of iterations is reached.